

COMP4680/8650 Advanced Topics in Statistical Machine Learning

Week 6: Dynamic Bayesian Networks

Stephen Gould

Temporal Models

In many situations we are interested in reasoning about the state of the world as it evolves over time. Examples:

- robot localization, patient health, speech recognition

Instead of a single set of random variables X we model the variables' trajectory over time $X_{0:T} = X_0, \dots, X_T$.


We will make two simplifications:

- time discretization
- Markov assumption

Markov Models

Markov Assumption: “the future is independent of the past given that we know the present”

$$p(X_{0:T}) = p_0(X_0) \prod_{t=1}^T p_t(X_t \mid X_{t-1})$$



Stationary Dynamical Systems

We say that a dynamical system is **stationary** (or time invariant) if $p_t(X_t \mid X_{t-1})$ is the same for all t .


If this is the case then we can represent infinite trajectories very compactly. We only need to represent

- the initial state distribution $p_0(X_0)$
- the transition model $p(X' \mid X)$


Dynamic Bayesian Networks

A **dynamic Bayesian network (DBN)** is a pair (B_0, B_{\rightarrow}) where

- B_0 is a BN over the variables at time zero
- B_{\rightarrow} is a 2-time-slice “template” defining
 - the dependence of variables from time $t-1$ to time t
 - the dependence of variables within time t



Typical DBN (Hidden Markov Model)



Inference in DBNs

- Consider a system evolving over time where:
 - x_t is the (unknown) system state at time t
 - u_t is the (known) control input at time t
 - y_t is some (observed) measurement at time t
- **filtering** estimates $p(x_{0:t} \mid y_{1:t}, u_{1:t})$
- **smoothing** estimates $p(x_{0:t} \mid y_{1:T}, u_{1:T})$
 - smoothing takes into account future observations

Efficient Recursive Filtering

Define the belief over the state at time t as

$$bel(x_t) = p(x_t \mid y_{1:t}, u_{1:t})$$

Bayesian updating proceeds in two steps:

- **prediction:** $bel'(x_t) = \int p(x_t \mid x_{t-1}, u_t) bel(x_{t-1}) dx_{t-1}$
- **correction:** $bel(x_t) \propto p(y_t \mid x_t) bel'(x_t)$

In order to implement the Bayes filter we often need to use tractable approximations. Considerations:

- computational efficiency, accuracy of approximation, ease of implementation

Linear Gaussian Systems

The linear Gaussian assumption for the **transition** (motion) and **measurement** models is a popular choice and leads to the famous Kalman filter.

The variables of the model are all assumed to be continuous-valued vectors:

- state (hidden), $x_t \in \mathbf{R}^n$
- control (input), $u_t \in \mathbf{R}^q$
- measurement (output), $y_t \in \mathbf{R}^p$

Linear Gaussian Assumptions

- The **state** is assumed to evolve as a linear function of the previous state and the current control input,

$$x_t = A_t x_{t-1} + B_t u_t + \eta_t$$

where ϵ_t is additive zero-mean Gaussian noise.

- The **measurements** are also a linear function of state,

$$y_t = C_t x_t + \delta_t$$

where δ_t is additive zero-mean Gaussian noise.

- Our belief over the **initial state**, x_0 , must also be Gaussian distributed.

Linear Gaussian Probabilities

- Initial belief: $p(x_0) = \mathcal{N}(x_0; \mu_0, \Sigma_0)$
- State transition probability:
$$p(x_t \mid u_t, x_{t-1}) = \mathcal{N}(x_t; A_t x_{t-1} + B_t u_t, Q_t)$$
- Measurement probability:
$$p(y_t \mid x_t) = \mathcal{N}(y_t; C_t x_t, R_t)$$

Under these assumptions the posterior distribution over x_t will be Gaussian.

Kalman Filter Algorithm

Prediction Step


$$\begin{aligned}\mu_{t|t-1} &= A_t \mu_{t-1|t-1} + B_t u_t \\ \Sigma_{t|t-1} &= A_t \Sigma_{t-1|t-1} A_t^T + Q_t\end{aligned}$$

Correction Step

$$\begin{aligned}K_t &= \Sigma_{t|t-1} C_t^T (C_t \Sigma_{t|t-1} C_t^T + R_t)^{-1} \\ \mu_{t|t} &= \mu_{t|t-1} + K_t (y_t - C_t \mu_{t|t-1}) \\ \Sigma_{t|t} &= (I - K_t C_t) \Sigma_{t|t-1}\end{aligned}$$

Computational efficiency: $O(p^3)$ for matrix inversion and $O(n^2)$ for covariance update. Measurement does not appear in Σ update.

Kalman Filter Example



Conditional Gaussians

Consider jointly Gaussian random variables x and y ,

$$\begin{bmatrix} x \\ y \end{bmatrix} \sim N \left(\begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix} \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{xy}^T & \Sigma_{yy} \end{bmatrix} \right)$$

Then $P(x | y)$ is Gaussian with

$$\begin{aligned} \mu_{x|y} &= \mu_x + \Sigma_{xy} \Sigma_{yy}^{-1} (y - \mu_y) \\ \Sigma_{x|y} &= \Sigma_{xx} - \Sigma_{xy} \Sigma_{yy}^{-1} \Sigma_{xy}^T \end{aligned}$$

(See Murphy §4.3.4.3)

Kalman Filter Derivation

Prediction Step

$$\begin{aligned}\mu_{t|t-1} &= E[x_t] \\ &= E[Ax_{t-1} + Bu_t + \eta_t] \\ &= AE[x_{t-1}] + Bu_t + E[\eta_t] \\ &= A\mu_{t-1|t-1} + Bu_t\end{aligned}$$

$$\begin{aligned}\Sigma_{t|t-1} &= \text{Var}[x_t] \\ &= \text{Var}[Ax_{t-1} + Bu_t + \eta_t] \\ &= A\text{Var}[x_{t-1}]A^T + \text{Var}[\eta_t] \\ &= A\Sigma_{t-1|t-1}A^T + Q_t\end{aligned}$$

Kalman Filter Derivation

Correction Step. First write as joint Gaussian over state and measurements

$$\begin{bmatrix} x_t \\ y_t \end{bmatrix} \sim N \left(\begin{bmatrix} \mu_{t|t-1} \\ C_t \mu_{t|t-1} \end{bmatrix} \begin{bmatrix} \Sigma_{t|t-1} & \Sigma_{t|t-1} C_t^T \\ C_t \Sigma_{t|t-1}^T & C_t \Sigma_{t|t-1} C_t^T + R_t \end{bmatrix} \right)$$

Now apply the conditional Gaussian rule,

$$\begin{aligned} \mu_{t|t} &= \mu_{t|t-1} + \Sigma_{t|t-1} C_t^T (C_t \Sigma_{t|t-1} C_t^T + R_t)^{-1} (y_t - C_t \mu_{t|t-1}) \\ \Sigma_{t|t} &= \Sigma_{t|t-1} - \Sigma_{t|t-1} C_t^T (C_t \Sigma_{t|t-1} C_t^T + R_t)^{-1} C_t \Sigma_{t|t-1}^T \end{aligned}$$

Extended Kalman Filter

The **extended Kalman filter** (EKF) allows the use of non-linear motion and measurement models,

$$\begin{aligned}x_t &= g(x_{t-1}, u_t) + \eta_t \\ y_t &= h(x_t) + \delta_t\end{aligned}$$

These models are linearized around the current mean of the state posterior, $\mu_{t-1|t-1}$, before each filter update

$$\begin{aligned}g(x_{t-1}, u_t) &\approx g(\mu_{t-1|t-1}, u_t) + G_t(x_{t-1} - \mu_{t-1|t-1}) \\ h(x_t) &\approx h(\mu_{t|t-1}) + H_t(x_t - \mu_{t|t-1})\end{aligned}$$

where G_t and H_t are Jacobians.

The accuracy of the EKF depends heavily on the goodness of the linear approximation.

Extended Kalman Filter

The **extended Kalman filter** (EKF) allows the use of non-linear motion and measurement models,

$$\begin{aligned}x_t &= g(x_{t-1}, u_t) + \eta_t \\ y_t &= h(x_t) + \delta_t\end{aligned}$$

These models are linearized around the current mean of the state posterior, $\mu_{t-1|t-1}$, before each filter update

$$\begin{aligned}g(x_{t-1}, u_t) &\approx g(\mu_{t-1|t-1}, u_t) + G_t(x_{t-1} - \mu_{t-1|t-1}) \\ h(x_t) &\approx h(\mu_{t|t-1}) + H_t(x_t - \mu_{t|t-1})\end{aligned}$$




where G_t and H_t are Jacobians.

constant

linear fcn of x_t

The accuracy of the EKF depends heavily on the goodness of the linear approximation.

Extended Kalman Filter Illustration



Extended Kalman Filter Algorithm

The extended Kalman filter is then

$$\begin{aligned}\mu_{t|t-1} &= g(\mu_{t-1|t-1}, u_t) \\ \Sigma_{t|t-1} &= G_t \Sigma_{t-1|t-1} G_t^T + R_t\end{aligned}$$

$$\begin{aligned}K_t &= \Sigma_{t|t-1} H_t^T (H_t \Sigma_{t|t-1} H_t^T + Q_t)^{-1} \\ \mu_{t|t} &= \mu_{t|t-1} + K_t (y_t - h(\mu_{t|t-1})) \\ \Sigma_{t|t} &= (I - K_t H_t) \Sigma_{t|t-1}\end{aligned}$$

Unscented Kalman Filter

Instead of using a Taylor series expansion to linearize the motion and measurement models, the **unscented Kalman filter** (UKF) performs stochastic linearization, which preserves statistical properties of the belief state.

- Weighted samples are generated from $N(\mu_{t-1|t-1}, \Sigma_{t-1|t-1})$ and propagated through $g(x_{t-1}, u_t)$
- A Gaussian is then fitted to the transformed results

The UKF does not require calculation of analytical derivatives. It also works better in practice than the EKF.

Particle Filtering

- The particle filter is a nonparametric Bayes filter, also known as **Sequential Monte Carlo**.
- The posterior $p(x_t \mid x_{0:t-1}, y_{1:t})$ is represented by a set of random samples

$$\mathcal{X}_t = \{x_t^{(1)}, \dots, x_t^{(m)}\}$$

- The representation is approximate (unlike the Kalman filter) but can model a much broader range of probability distributions (and non-linear dynamics).

Particle Filter Algorithm

Update and Weight Particles

for $i = 1$ to m do

sample $x_t^{(i)}$ from $p(x_t | u_t, x_{t-1}^{(i)})$

set $w_t^{(i)} = p(y_t | x_t^{(i)})$

Resample Particles


initialise $\mathcal{X}_t = \emptyset$

for $j = 1$ to m do


draw i with probability proportional to $w_t^{(k)}$

add $x_t^{(i)}$ to \mathcal{X}_t

Particle Filter Example




real robot location




particle filter

Resampling Step


**unweighted
particles**



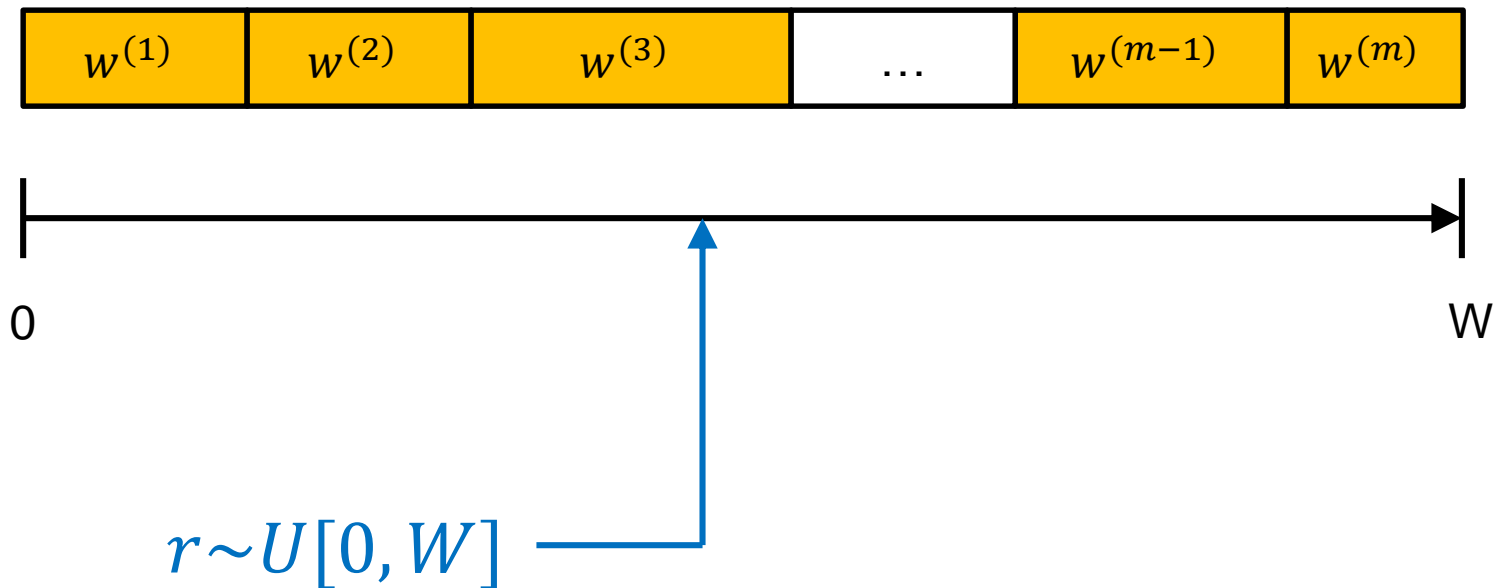
**weighted
particles**




**unweighted
resampled
particles**



Sampling from $[0, W]$



Low Variance Sampler



$$r \sim U[0, m^{-1}W] \quad r + m^{-1}W \quad r + km^{-1}W$$

makes samples dependent

Low Variance Sampler Algorithm

initialise $\mathcal{X}_t = \emptyset$

set $W = \sum_{i=1}^m w_t^{(i)}$

sample $r \sim U[0, m^{-1}W]$

set $i = 1$ and $c = w_t^{(1)}$

for $j = 1$ to m do

$U = r + (j - 1)m^{-1}W$

 while $U > c$ do

 set $i += 1$ and $c += w_t^{(i)}$

 add x_t^i to \mathcal{X}_t

Particle Deprivation

- **Particle Deprivation:** no particles are near the true state
- This is a result of variance in random sampling
 - An unlucky series of random numbers can wipe out all particles near the true state
 - Since this has a non-zero probability of happening at each time slice it will happen eventually
- A popular solution is to add a small number of randomly generated particles when resampling
 - Also addresses the “kidnapped robot” problem

Distributional Particles

- Computational cost/number of particles explodes as dimensionality increases
- Distributional particles (aka **Rao-Blackwellization**) by keeping part of the distribution in analytical form
- Example: conditionally linear Gaussian models
 - Let the state be represented by $(x_t, z_t) \in \mathbf{R}^n \times [1, \dots, K]$
 - Assume that conditioned on z_t the system can be modelled with linear Gaussian dynamics
$$x_t = A(z_{t-1})x_{t-1} + B(z_{t-1})u_t + \eta_t$$
$$y_t = C(z_t)x_t + \eta_t$$
 - Then we can approximate the distribution over (x_t, z_t) by a set of particles with $z_t \sim [0, \dots, K]$ and $x_t \sim N(\mu_t, \Sigma_t)$.

RBPF Algorithm

Update and Weight Particles

for $i = 1$ to m do

sample $z_t^{(i)}$ from $p(z_t | u_t, z_{t-1}^{(i)})$

Kalman filter update for $\mu_t^{(i)}$ and $\Sigma_t^{(i)}$ given $z_t^{(i)}$

set $w_t^{(i)} = p(y_t | x_t^{(i)}, z_t^{(i)})$

Resample Particles

initialise $\mathcal{X}_t = \emptyset$

for $j = 1$ to m do

draw i with probability proportional to $w_t^{(k)}$

add $(x_t^{(i)}, z_t^{(i)})$ to \mathcal{X}_t