

# STRUCTURAL SVMS

---

COMP 4680/8650

Justin Domke

[comp8650@anu.edu.au](mailto:comp8650@anu.edu.au)

Reading: Murphy, §19.7

# What to do for inference?

- Suppose we have some distribution  $p(y|x)$ .
  - E.g.  $x$  = images,  $y$  = segmentations.
- Given an input  $x$ , what to predict?

$$y^* = \arg \max_{y \in \mathcal{Y}} p(y|x)$$

$$y_i^* = \arg \max_{y_i \in \mathcal{Y}_i} p(y_i|x)$$

- What's the right thing to do?

# Utility Functions

- $U(y, a)$  says how happy you are to take action  $a$  if true output is  $y$ .
- E.g., driving a car.

$$a \in \{\text{left, center, right}\}$$

$$U(y, a) = \begin{cases} 1 & \text{if on road} \\ -10 & \text{if off road} \\ -100000 & \text{if crashed} \end{cases}$$

# Utility Functions

- Pick action  $a$  to maximize expected reward.

$$\sum_y p(y|x)U(y, a)$$

- Sometimes “actions” are the states themselves.

- If  $U(y, y') = I[y = y']$ , then optimal approach is

$$y' = \arg \max_{y \in \mathcal{Y}} p(y|x)$$

- If  $U(y, y') = \sum_i I[y_i = y'_i]$ , then optimal approach is

$$y'_i = \arg \max_{y_i \in \mathcal{Y}_i} p(y_i|x)$$

# MAP Inference

- If we are only going to use  $p(y|x; w)$  to do MAP inference, why fit a full probability distribution?
- Alternative Approach. Define the “mapping”

$$y^*(x; w) := \arg \max_{y \in \mathcal{Y}} p(y|x; w)$$

and just find  $w$  to make this mapping accurate.

- Why do this?
  - Model mis-specification. (Similar to conditional vs. joint training.)
  - Computational Issues.

# Computational Issues

- MAP inference is NP-complete
- Marginal inference is #P-complete
- Tons of works exists on great MAP solvers, often for specific problems (proteins, image segmentations)
- Often, MAP solvers can certify that they found the optimal solution in a given instance, even if this is not possible in general.

# SSVMs – Basic idea

- Take some probabilistic model

$$p(y|x; w) = \exp(w^T \phi(y, x) - A(w))$$

- Instead of being probabilistic, define a loss directly on the inference procedure

$$y^*(x; w) := \arg \max_{y \in \mathcal{Y}} p(y|x; w)$$

$$y^*(x; w) = \arg \max_{y \in \mathcal{Y}} w^T \phi(y, x)$$

# Structured Perceptron

$$\min_w \sum_i Q(w, x_i, y_i)$$

$$Q(w, x_i, y_i) = \max_{\hat{y}_i \in \mathcal{Y}} w^T \phi(\hat{y}_i, x_i) - w^T \phi(y_i, x_i)$$

- What needs to happen to have no loss?

$$w^T \phi(y_i, x_i) = \max_{\hat{y}_i \in \mathcal{Y}} w^T \phi(\hat{y}_i, x_i)$$



# Structured Perceptron

## Stochastic subgradient descent

- For  $k = 1, 2, \dots$ 
  - Pick a random  $i$
  - Find  $\hat{y}_i = \arg \max_{\hat{y}_i \in \mathcal{Y}} w^T \phi(\hat{y}_i, x_i)$
  - Set  $w \leftarrow w - \eta_k \cdot (\phi(\hat{y}_i, x_i) - \phi(y_i, x_i))$

Gradually decrease  $\eta_k$  over time.

# Structured Perceptron

- Advantages:
  - Only need to be able to perform MAP inference
- Disadvantages
  - Only need score of correct output to slightly beat all incorrect output
  - Loss is not tuned for application priorities.

# Structured SVM

$$\min_w \sum_i Q(w, x_i, y_i)$$

$$Q(w, x_i, y_i) = \max_{\hat{y}_i \in \mathcal{Y}} L(y_i, \hat{y}_i) + w^T \phi(\hat{y}_i, x_i) - w^T \phi(y_i, x_i)$$

- What needs to happen to have no loss?

$$w^T \phi(y_i, x_i) = \max_{\hat{y}_i \in \mathcal{Y}} L(y_i, \hat{y}_i) + w^T \phi(\hat{y}_i, x_i)$$

# Structural SVMs

- $L(y_i, \hat{y}_i)$  measures how much you dislike predicting  $\hat{y}_i$  when true output is  $y_i$ .
- Common example: Hamming distance

$$L(y_i, \hat{y}_i) = \sum_n I[y_{in} \neq \hat{y}_{in}]$$

# Structured SVMs

## Stochastic subgradient descent

- For  $k = 1, 2, \dots$ 
  - Pick a random  $i$
  - Find  $\hat{y}_i = \arg \max_{\hat{y}_i \in \mathcal{Y}} L(y_i, \hat{y}_i) + w^T \phi(\hat{y}_i, x_i)$
  - Set  $w \leftarrow w - \eta_k \cdot (\phi(\hat{y}_i, x_i) - \phi(y_i, x_i))$

Gradually decrease  $\eta_k$  over time.

Loss-augmented inference



# Loss-Augmented Inference

$$\hat{y}_i = \arg \max_{\hat{y}_i \in \mathcal{Y}} L(y_i, \hat{y}_i) + w^T \phi(\hat{y}_i, x_i)$$

- Typically,  $L$  can be “folded into” into the features, so that

$$\hat{y}_i = \arg \max_{\hat{y}_i \in \mathcal{Y}} w'^T \phi(\hat{y}_i, x_i)$$

- Example:

$$\phi(y, x) = \{x \cdot I[y_i = y_i^*]\} \cup \{I[y_i = y_i^*, y_j = y_j^*]\}$$


$$L(y_i, \hat{y}_i) = \sum_i I[y_i \neq \hat{y}_i]$$

# Demo – Multi-Label Prediction

$$x \in \mathbb{R}^N \quad y \in \{-1, +1\}^N \quad L(y_i, \hat{y}_i) = \sum_i I[y_i \neq \hat{y}_i]$$



With two variables:  $\phi(y, x) = (x_1 y_1, x_2 y_2, y_1 y_2)$



With N variables:  $\phi(y, x) = \{x_i y_i\} \cup \{y_i y_{i+1}\}$

# Optimization Theory

$$\min_w \frac{\lambda}{2} ||w||^2 + \sum_{i=1}^N \max_y w^T (\phi(y, x_i) - \phi(y_i, x_i)) + L(y_i, y)$$

$$\min_w \frac{1}{2} ||w||^2 + C \sum_{i=1}^N \max_y w^T (\phi(y, x_i) - \phi(y_i, x_i)) + L(y_i, y)$$

$$\min_w \frac{1}{2} ||w||^2 + C \sum_{i=1}^N \xi_i$$

$$\text{s.t. } \forall i, \forall y, \xi_i \geq L(y_i, y) + w^T (\phi(y, x_i) - \phi(y_i, x_i))$$



# Optimization Theory

$$\begin{aligned} \min_{w, \xi} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & \forall i, \forall y, \xi_i \geq L(y_i, y) + w^T (\phi(y, x_i) - \phi(y_i, x_i)) \end{aligned}$$

- This is a quadratic objective, with linear constraints, known as a quadratic program (QP).
- However, there are  $N|\mathcal{Y}|$  constraints!
  - E.g., with binary variables,  $|\mathcal{Y}| = 2^N$ .

# Max-Margin Markov Networks (2003)

- Convert optimization to “dual form”.

**Primal formulation (6)**

**Dual formulation (7)**

$$\begin{array}{ll} \min & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{\mathbf{x}} \xi_{\mathbf{x}} ; \\ \text{s.t.} & \mathbf{w}^\top \Delta \mathbf{f}_{\mathbf{x}}(\mathbf{y}) \geq \Delta \mathbf{t}_{\mathbf{x}}(\mathbf{y}) - \xi_{\mathbf{x}}, \forall \mathbf{x}, \mathbf{y}. \end{array} \quad \begin{array}{ll} \max & \sum_{\mathbf{x}, \mathbf{y}} \alpha_{\mathbf{x}}(\mathbf{y}) \Delta \mathbf{t}_{\mathbf{x}}(\mathbf{y}) - \frac{1}{2} \left\| \sum_{\mathbf{x}, \mathbf{y}} \alpha_{\mathbf{x}}(\mathbf{y}) \Delta \mathbf{f}_{\mathbf{x}}(\mathbf{y}) \right\|^2 ; \\ \text{s.t.} & \sum_{\mathbf{y}} \alpha_{\mathbf{x}}(\mathbf{y}) = C, \forall \mathbf{x}; \quad \alpha_{\mathbf{x}}(\mathbf{y}) \geq 0, \forall \mathbf{x}, \mathbf{y}. \end{array}$$

- Make change of variables, optimize over marginals.

$$\begin{aligned} \mu_{\mathbf{x}}(y_i, y_j) &= \sum_{\mathbf{y} \sim [y_i, y_j]} \alpha_{\mathbf{x}}(\mathbf{y}), & \forall (i, j) \in E, \forall y_i, y_j, \forall \mathbf{x}; \\ \mu_{\mathbf{x}}(y_i) &= \sum_{\mathbf{y} \sim [y_i]} \alpha_{\mathbf{x}}(\mathbf{y}), & \forall i, \forall y_i, \forall \mathbf{x}; \end{aligned}$$

(Only works for treelike graphs.)

# Cutting Plane Optimization (2009)

- Repeat: 1) Solve this optimization:


$$\begin{aligned} \min_{w, \xi} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & \forall (i, y) \in A, \quad \xi_i \geq L(y_i, y) + w^T (\phi(y, x_i) - \phi(y_i, x_i)) \end{aligned}$$

- 2) For all  $i$ , find most violated constraint:

$$\max_y L(y_i, y) + w^T (\phi(y, x_i) - \phi(y_i, x_i))$$

- If more than  $\xi_i$ , add  $(i, y)$  to  $A$ .
- 3) If don't need to add any constraints, we have the global optimum!

# Cutting Plane Optimization (2009)



# Linear Complexity: SVM-Struct (2009)

$$\begin{aligned} \min_{w, \xi} \quad & \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & \forall (i, y) \in A, \quad \xi_i \geq L(y_i, y) + w^T (\phi(y, x_i) - \phi(y_i, x_i)) \end{aligned}$$

- Reformulate with a single  $\xi$

$$\begin{aligned} \min_{w, \xi} \quad & \frac{1}{2} \|w\|^2 + C\xi \\ \text{s.t.} \quad & \forall (\bar{y}_1, \dots, \bar{y}_N) \in A, \quad \xi \geq \sum_{i=1}^N L(y_i, \bar{y}_i) + w^T (\phi(\bar{y}_i, x_i) - \phi(y_i, x_i)) \end{aligned}$$

- Same strategy: iteratively add to A. But, can show:
  - Optimizing objective takes  $O(N)$  time.
  - Total number of iterations is a constant.
  - Thus, solve entire optimization in  $O(N)$  time, and  $O(N)$  oracle calls.